



ECON6027

1c

Handling Vector Spatial Data



Sf objects



In R, sf objects are datasets that have spatial awareness. That means they know exactly “where” we are talking about!



In 1b we saw how to import a geographic dataset to R and thereby create a sf object.



In this section, we will discuss explore sf objects further.



Packages you need

- sf
- spData
 - contains spatial datasets

Import data to r

- Import the csv file to the R using read.csv() function.
- We have a spatial dataset, but this dataset has no "spatial awareness"
- Let's give it some spatial powers!

Give spatial awareness

“coords” argument follows the cartesian coordinates convention (x,y).

```
> fav.sf = st_as_sf(fav, coords=c("lon_x", "lat_y"))
> fav.sf
Simple feature collection with 9 features and 3 fields
Geometry type: POINT
Dimension: XY
Bounding box: xmin: -139 ymin: -66 xmax: 118 ymax: 85
CRS: NA
```

	Name	num	col	geometry
1	ABC	10	red	POINT (-134 -54)
2	DEF	3	yellow	POINT (118 67)
3	GHI	5	blue	POINT (-17 67)
4	JKL	10	green	POINT (83 85)
5	MNO	5	purple	POINT (-123 21)
6	PQR	4	black	POINT (18 -46)
7	STU	1	white	POINT (-139 -66)
8	VW	5	pink	POINT (-4 29)
9	XYZ	8	orange	POINT (-130 -53)


Basic Terminology

- One sfg object contains only a single simple feature geometry.
- A simple feature geometry column (sfc) is a list of sfg objects, which is additionally able to contain information about the coordinate reference system in use. (Ch 2)
- This is important since sfc represents the geometry column in **sf** data frames

```
> fav.sf = st_as_sf(fav, coords=c("lon_x", "lat_y"))
> fav.sf
```

Simple feature collection with 9 features and 3 fields
Geometry type: POINT
Dimension: XY
Bounding box: xmin: -139 ymin: -66 xmax: 118 ymax: 85
CRS: NA

	Name	num	col	geometry
1	ABC	10	red	POINT (-134 -54)
2	DEF	3	yellow	POINT (118 67)
3	GHI	5	blue	POINT (-17 67)
4	JKL	10	green	POINT (83 85)
5	MNO	5	purple	POINT (-123 21)
6	PQR	4	black	POINT (18 -46)
7	STU	1	white	POINT (-139 -66)
8	VW	5	pink	POINT (-4 29)
9	XYZ	8	orange	POINT (-130 -53)



Roll up your sleeves...

Install package "spData" and load the "world" dataset as an sf object:

```
> world = st_read(system.file("shapes/world.gpkg",  
package="spData"))
```

Inspect the **world** dataset and general sf properties

```
> class(world); dim(world); names(world);  
summary(world); head(world)
```

sf class dataset

```
> class(world)
[1] "sf"          "tbl_df"      "tbl"         "data.frame"
```

Test yourself...

- Identify "attributes" and "features".
- sf objects also have a special column to contain geometry data, can you identify it?
- What is the coordinate reference system associated with this dataset?

```
> head(world)
```

Simple feature collection with 6 features and 10 fields

geometry type: MULTIPOLYGON

dimension: XY


bbox: xmin: -180 ymin: -18.28799 xmax: 180 ymax: 83.23324

CRS: EPSG:4326

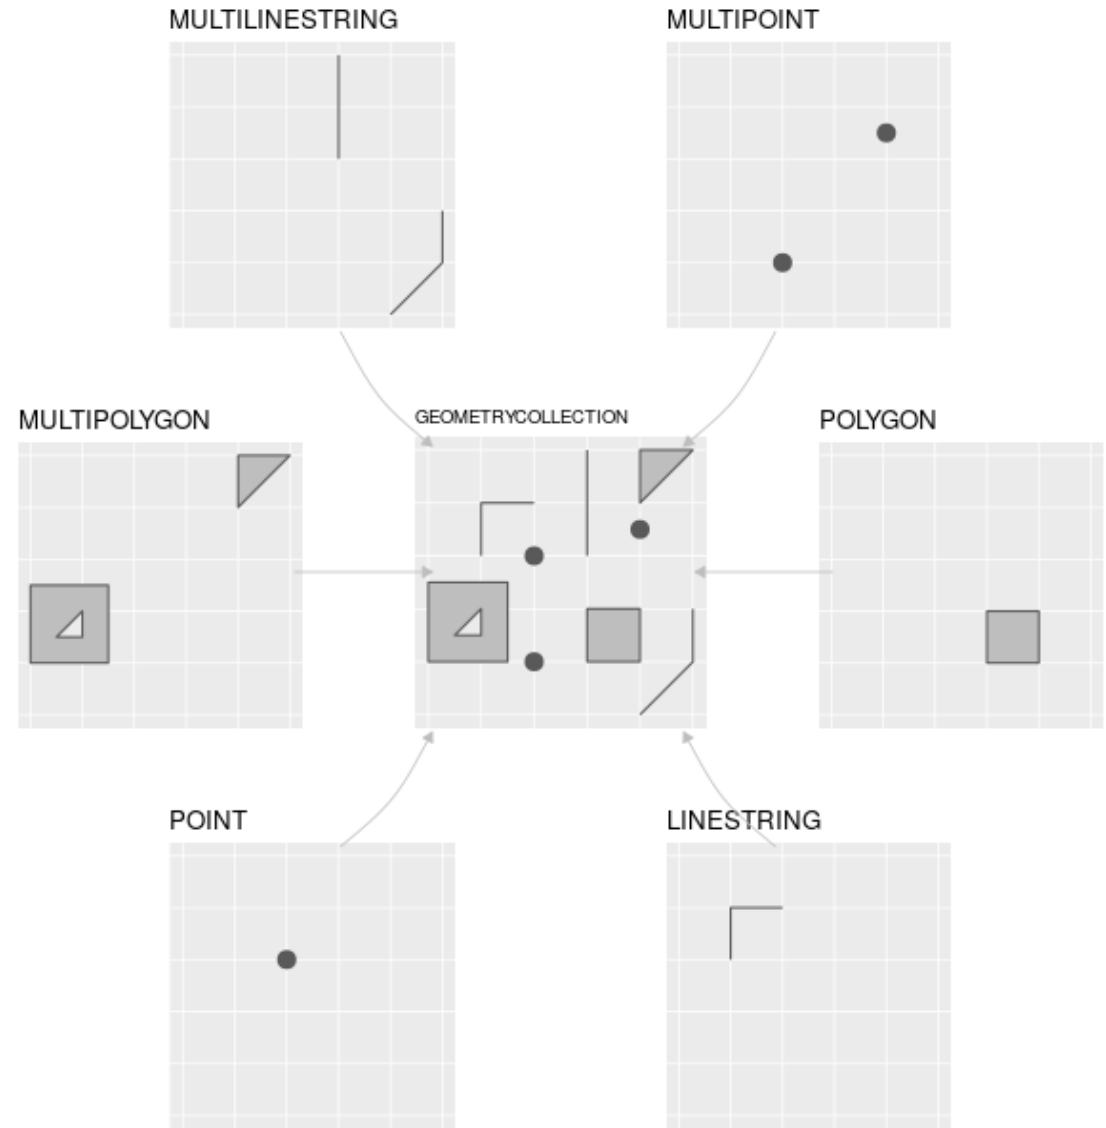
	iso_a2	name_long	continent	region_un	subregion	type	area_km2	pop	lifeExp	gdpPercap	geom
1	FJ	Fiji	Oceania	Oceania	Melanesia	Sovereign country	19289.97	885806	69.96000	8222.254	MULTIPOLYGON (((180 -16.067...
2	TZ	Tanzania	Africa	Africa	Eastern Africa	Sovereign country	932745.79	52234869	64.16300	2402.099	MULTIPOLYGON (((33.90371 -0...
3	EH	Western Sahara	Africa	Africa	Northern Africa	Indeterminate	96270.60	NA	NA	NA	MULTIPOLYGON (((-8.66559 27...
4	CA	Canada	North America	Americas	Northern America	Sovereign country	10036042.98	35535348	81.95305	43079.143	MULTIPOLYGON (((-122.84 49,...
5	US	United States	North America	Americas	Northern America	Country	9510743.74	318622525	78.84146	51921.985	MULTIPOLYGON (((-122.84 49,...
6	KZ	Kazakhstan	Asia	Asia	Central Asia	Sovereign country	2729810.51	17288285	71.62000	23587.338	MULTIPOLYGON (((87.35997 49...

Attributes vs. Features

- Attribute data is “non-spatial” information associated with geographic (geometry) data. Examples:
 - A bus stop provides a simple example: its position would typically be represented by latitude and longitude coordinates (geometry data), in addition to its name. The name is an *attribute* of the feature (to use Simple Features terminology) that bears no relation to its geometry.
 - The elevation value (attribute) for a specific grid cell in raster data.



Feature types supported by sf



Geometry column

The geometry column gives the sf object its “spatial awareness”.

It is a list column that contains all the coordinates of the country polygons.

Inspect the geometry features

```
> world$geom
```

```
Geometry set for 177 features
```


```
Geometry type: MULTIPOLYGON
```

```
Dimension:      XY
```

```
Bounding box:  xmin: -180 ymin: -89.9  
xmax: 180 ymax: 83.64513
```

```
Geodetic CRS:  WGS 84
```

```
First 5 geometries...
```

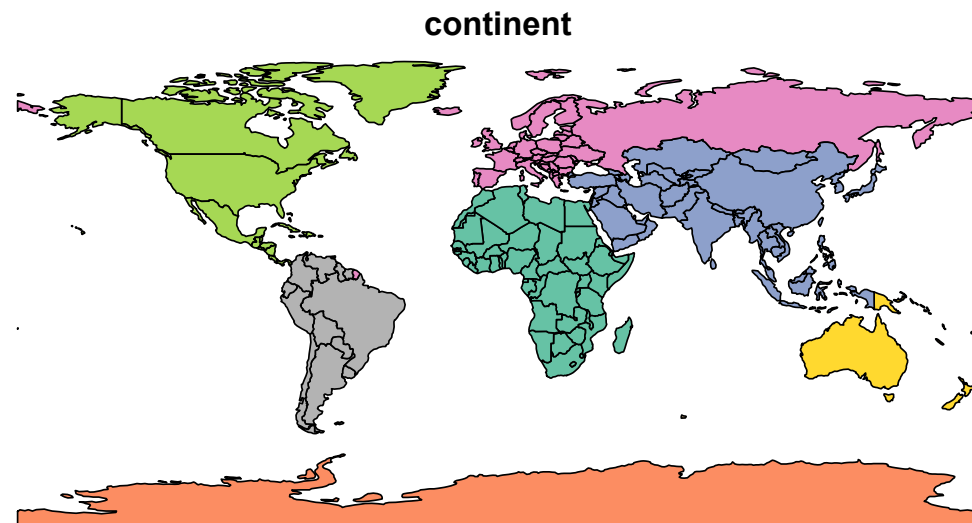
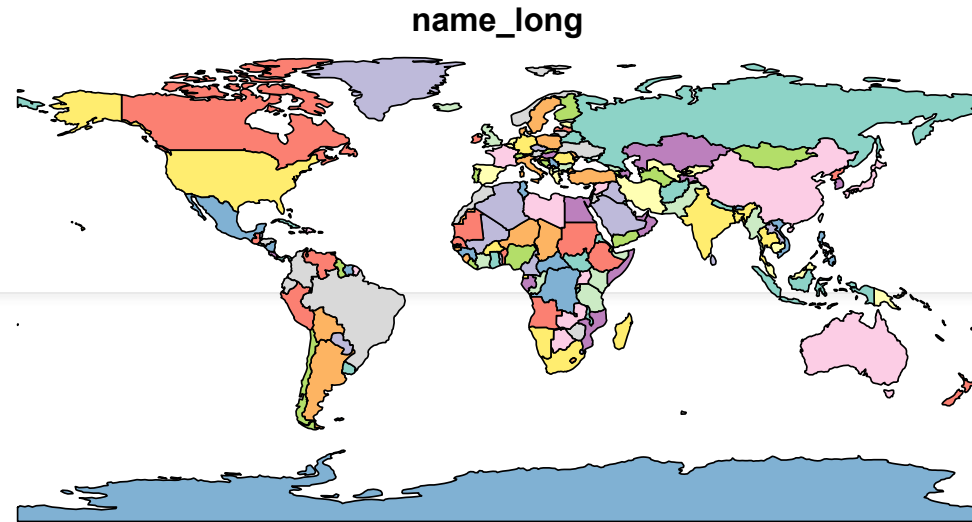



If you strip the geometry data from an "sf" object, it reduces to a typical "data.frame"

```
> world_df =  
st_drop_geometry(world);  
class(world_df)  
  
[1] "data.frame"  
  
> head(world_df)
```

Visualise

```
> plot(world) # plot all  
> plot(world[c("name_long",  
"continent")])
```






Manipulation of a Vector Dataset

1. Spatial (Attribute) Subsetting

- Spatial subsetting is the process of taking a spatial object and returning a new object containing only features that *relate* in space to another object.
- In other words, we will collect attributes that have some shared common element.
 - eg: collection of mall names and size in the central region



1. Subsetting Examples

Extract a numerical summary of the numerical data in columns 7 through 10:

```
> summary(world[, 7:10])
```

A subset of large countries:

```
> l_area = world$area_km2 > 5000000
```

```
> (big_countries = world[l_area, ])
```

```
# or
```

```
> (big_countries = subset(world, area_km2 > 5000000))
```

Standard comparison operators

Symbol	Name
==	Is equal to
!=	Is not equal to
<,>	Is great/less than
<=, >=	Is greater/less than or equal to
&, , !	and, or, not

Subsetting Examples (contd.)

- Remove Antarctica

```
> (big_countries = big_countries[-7, ])
```

Simple feature collection with 6 features and 10 fields

geometry type: MULTIPOLYGON

dimension: XY

bbox: xmin: -180 ymin: -43.6346 xmax: 180 ymax: 83.23324

CRS: EPSG:4326

A tibble: 6 x 11

	iso_a2	name_long	continent	region_un	subregion	type	area_km2	pop	lifeExp	gdpPercap	geom
	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<MULTIPOLYGON [°]>
1	CA	Canada	North Ame...	Americas	Northern A...	Sove...	1.00e7	3.55e7	82.0	43079.	(((-122.84 49, -122.9742 49....
2	US	United St...	North Ame...	Americas	Northern A...	Coun...	9.51e6	3.19e8	78.8	51922.	(((-122.84 49, -120 49, -117...
3	RU	Russian F...	Europe	Europe	Eastern Eu...	Sove...	1.70e7	1.44e8	70.7	25285.	(((178.7253 71.0988, 180 71....
4	BR	Brazil	South Ame...	Americas	South Amer...	Sove...	8.51e6	2.04e8	75.0	15374.	(((-53.37366 -33.76838, -53....
5	AU	Australia	Oceania	Oceania	Australia ...	Coun...	7.69e6	2.35e7	82.3	43547.	(((147.6893 -40.80826, 148.2...
6	CN	China	Asia	Asia	Eastern As...	Coun...	9.41e6	1.36e9	75.9	12759.	(((109.4752 18.1977, 108.655...



EXERCISE A

1. Create a new sf object made of columns 3 to 6.
2. Create a new sf object by removing "iso_a2" and "region_un".
3. Filter the countries with "gdpPercap" more than \$50,000.
4. Try at home: Generate a subset of the "richest three" countries in "Europe".

Subsetting Examples (contd.)

```
> world$continent # check listed  
continents
```

```
> asia = world[world$continent ==  
"Asia", ]
```

extract Asia. Using [] is one way to subset a dataset. Notice the resulting dataset is also an sf object.

```
> plot(world["continent"], reset = F)  
# if the first object has more than one key reset  
must be set to False for "add=T" option to work in  
the next line.
```

```
> plot(asia, add = T, col = "black")
```

Where do I come from?

```
> world$name_long # list of countries  
> (SL = world[world$name_long == "Sri Lanka", ]) # subsetting  
> plot(st_geometry(asia), main="Asia")  
> plot(st_geometry(SL), col = "red", lwd = 3, add=T)
```

Notice [0] and plot(st_geometry()) both gives the same outcome: an outline of the geometry column.

Asia






Exercise B

- Repeat the steps above for a country of your choice. (Unfortunately, Singapore is not available in this dataset).

2. Attribute Aggregation

Aggregation operations summarise datasets by a 'grouping variable', typically an attribute column.



Attribute data aggregation example

Aggregate the population by continent:

1. Using aggregate function in "stats" package (invoked when using ~)

```
> (cont_pop = aggregate(pop ~ continent, FUN = sum, data = world, na.rm = T)); class(cont_pop) # output is a dataframe
```

2. Using aggregate function in "sf" package (invoked when using an sf object and a "by" argument is given)

```
> (cont_pop2 = aggregate(world['pop'], by = list(world$continent), FUN = sum, na.rm = T)); class(cont_pop2) # output is an sf object
```

na.rm=T allows R to exclude missing values in the calculation



EXERCISE C

1. Can you explain what the difference is between "cont_pop" and "cont_pop2"?
2. Run the command "?aggregate" and inspect the function arguments.
3. Create an object of the average population in each continent using "FUN=mean" while maintaining the geometry column of the object.



3. Attribute Joining

Joining data from different sources

This is useful when you want to “give” spatial awareness to a dataset.

Example: attribute joining using a shared “key” variable

Combine “world” and “coffee_data” found in spData package.

```
> library(spData)
```

```
> summary(coffee_data); class(coffee_data)
```


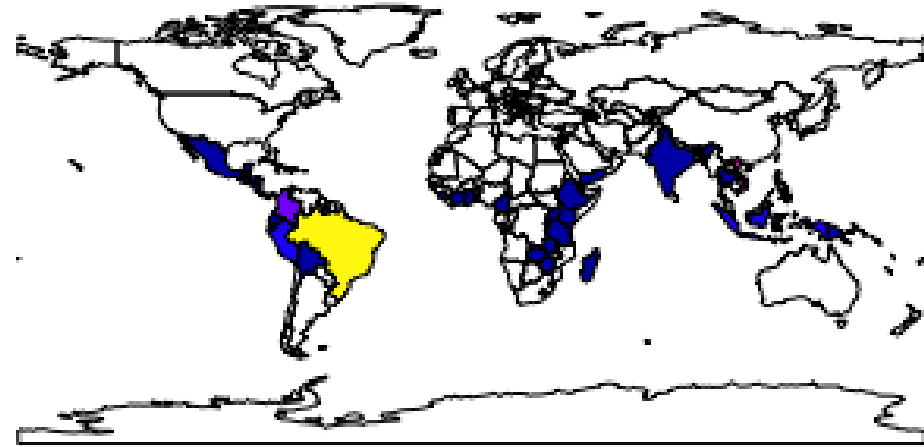
```
> (world_coffee = left_join(world, coffee_data))
```

left_join preserves the left/first dataset.

```
> ?left_join # see what other ways of joining two datasets are available.
```


Visualise the
added
variable

coffee_production_2016



The “key variable”

- Notice the joining is done using the shared “**key**” variable “name_long”.
- By default, **dplyr** uses all variables with matching names. In this case, both world_coffee and world objects contained a variable called name_long, explaining the message `Joining, by = "name_long"`.
- Where variable names are not the same, you have two options:
 - Rename the key variable in one of the objects so they match.
 - Use the `by` argument to specify the joining variables.



Other ways of mutating joins

- The mutating joins add columns from y to x, matching rows based on the keys:
 - `inner_join()`: includes all rows in x and y.
 - `left_join()`: includes all rows in x.
 - `right_join()`: includes all rows in y.
 - `full_join()`: includes all rows in x or y.
- If a row in x matches multiple rows in y, all the rows in y will be returned once for each matching row in x.

Example: attribute joining using a shared “key” variable

`left_join` preserves the left dataset in full. i.e., if there is an entry in the left dataset without an exact match with the right dataset, it simply inserts an `<NA>` entry. (Check entry 3 of `world_coffee`). What if we want to remove such entries completely. i.e. get an exact match of the left and right datasets?

```
> (world_coffee2 = inner_join(world, coffee_data))  
# the sf object must be in front if you want to  
retain the sf class
```



EXERCISE D

1. Illustrate coffee production in 2017 in a world map.
2. Conduct a `left_join` to create a dataset of only the coffee producing countries. Is this an sf object? How will you convert this to an sf object?



4. Create new attributes

```
> (world_coffee$prod_yoy =  
  (world_coffee$coffee_production_2017/world_c  
    offee$coffee_production_2016 - 1)*100)  
  
> world_coffee
```

Exercise E (Challenge Exercise)

1. Load the world dataset from spData. Generate a subset that consists of Asia. Create a preliminary plot.
2. Load the “arrivals.Rdata” and extract the arrivals from Asia. Give it spatial powers by performing an appropriate joining with the world dataset. The final object must have all the attributes from both datasets. Inspect the object.
3. Extract a subset of arrivals from 2000 onward.
4. Find the average arrivals from each Asian country between 2000 to 2018. Plot it.
5. Add a population column to the average arrivals. Create a new attribute that shows the annual average arrivals as a percentage of total population. Plot this attribute.

Take home points...

- Manipulate attributes and features of an sf object
 - Subsetting
 - Aggregation
 - Joining



References

- <https://geocompr.robinlovelace.net/attr.html#vector-attribute-manipulation>
- <https://geocompr.robinlovelace.net/spatial-operations.html#spatial-vec>
- Datasets: <https://nowosad.github.io/spData/>